

Parallel Programming in Python – a very, very short introduction

Agenda

1. Computer architectures, data structures, and opportunities
2. Case 1 – computing a distance matrix
3. Case 2 – computing the average of a large sample

The scope of the technique that will be presented [1]:

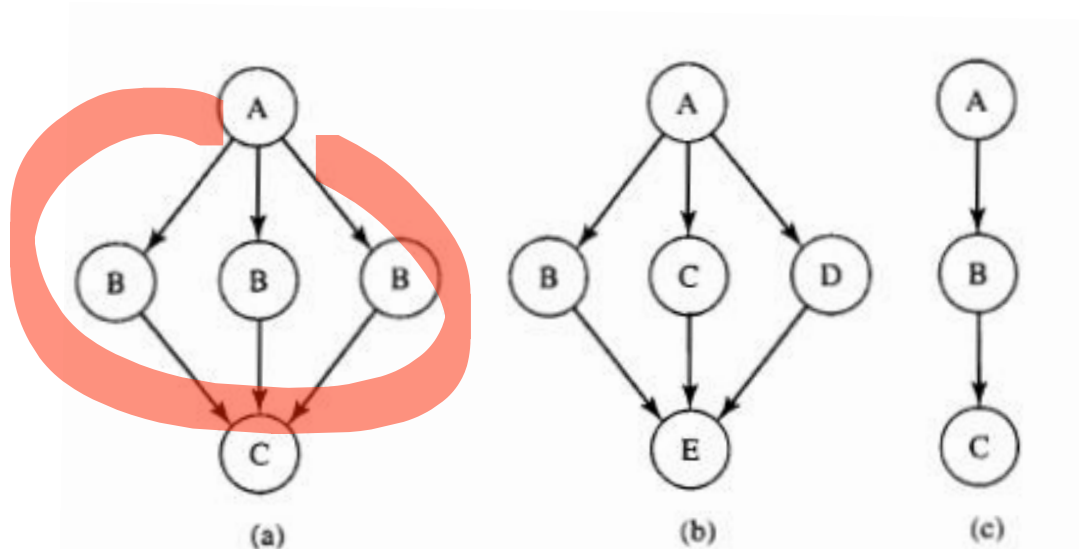


Figure 1.4 Parallelism in data dependence graphs. Vertices represent tasks. The letter inside a vertex indicates the operation being performed. Edges denote dependences among tasks. (a) A graph exhibiting data parallelism. Three tasks may concurrently apply operation B to different operands. (b) A graph exhibiting functional parallelism. Tasks performing operations B, C, and D may be performed concurrently. (c) A purely sequential dependence graph. However, if all tasks take the same amount of time to execute and multiple problem instances need to be processed, operation C may be performed on instance i while operation B is performed on instance $i + 1$ and operation A is performed on instance $i + 2$. This structure is called pipelining.

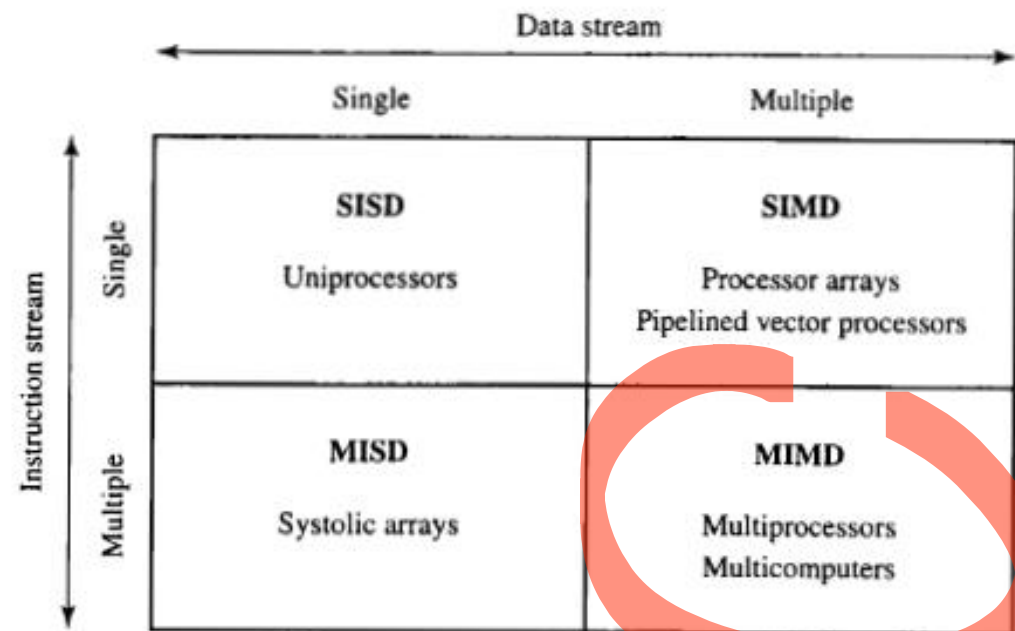


Figure 2.20 Flynn's taxonomy of computer architectures.

Resumo. *Aplicações Bag-of-Tasks (BoT) são aplicações paralelas compostas de tarefas independentes (ou seja, embaraçosamente paralelas), que não se comunicam entre si, podem depender de um ou mais arquivos de entrada e podem ser executadas em qualquer orden [2]*



Case 1 – computing a distance matrix

Case 1 - computing a distance matrix, sequential scheme



distance matrix

	v1	v2	v3		vk
v1	0				
v2		0			v
v3			0		
				0	
vk		v			0

tasks

v1	v1	v1	v1	v2	v2	v2	v3	v3	
v2	v3		vk	v3		vk		vk	vk



Sequential execution, process P1



--	--	--	--	--	--	--	--	--	--

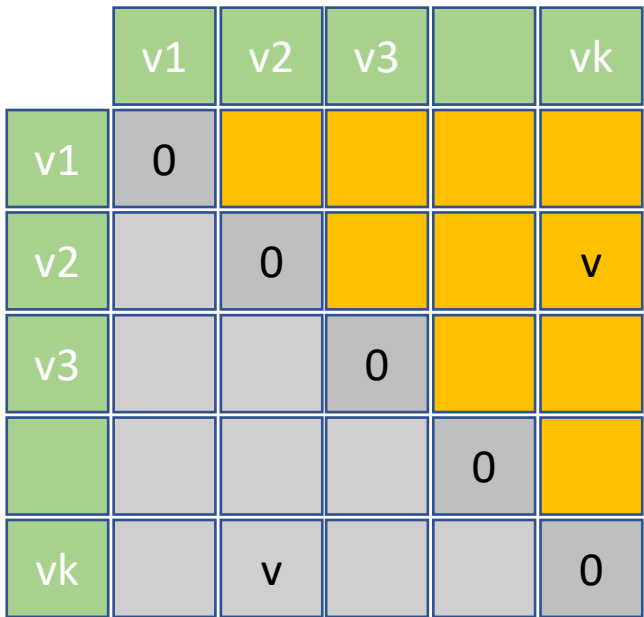
distance matrix

Case 1 - computing a distance matrix, parallel scheme


d-dimensional vectors



distance
matrix



tasks



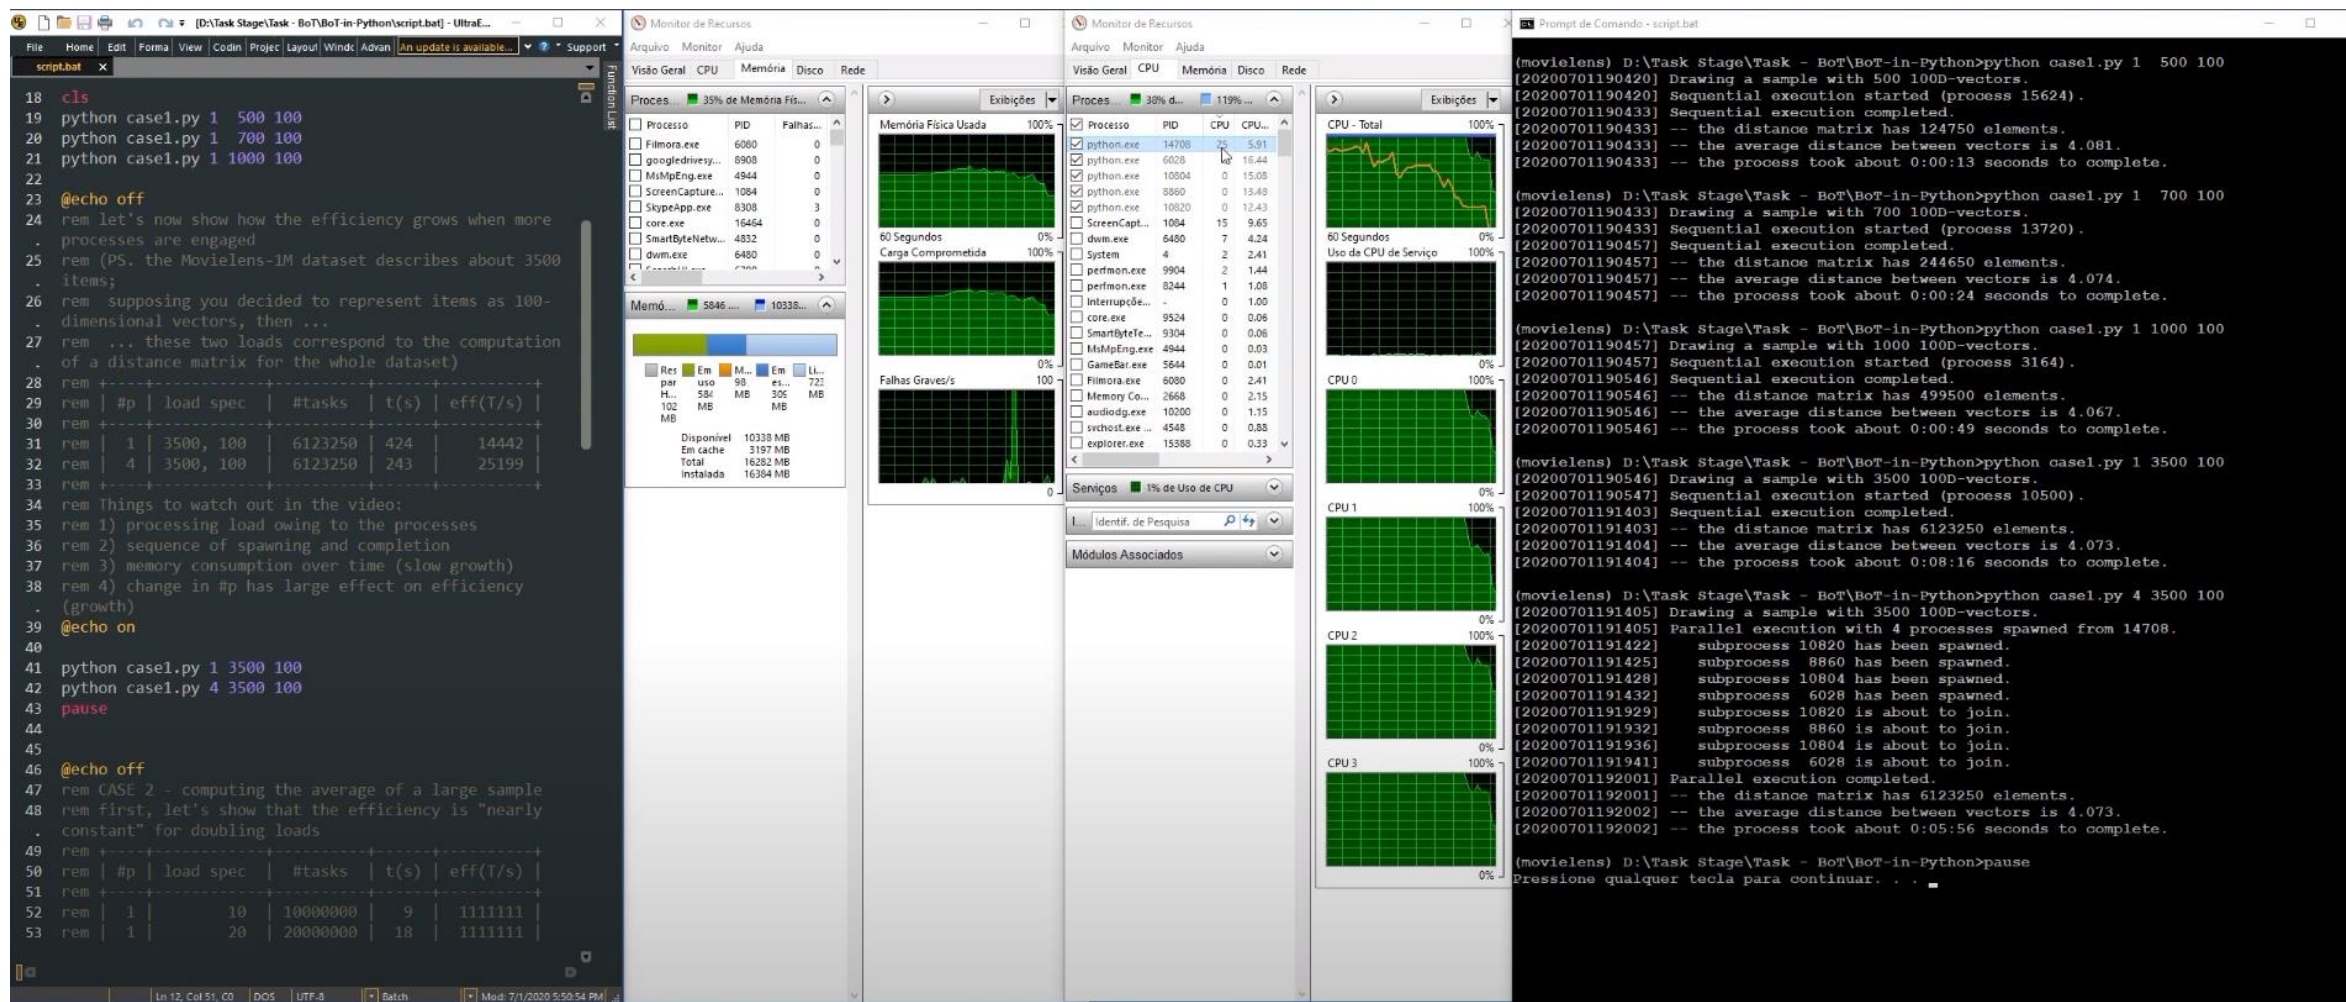
Parallel execution process P1

Parallel execution process P2



distance matrix

Case 1 - computing a distance matrix, sequential scheme



[Link para o vídeo](#)




Case 2 – computing the average of a large sample

Case 2 - computing the average of a large sample, sequential scheme

(sampled from the Brazilian male population, 1977 [3])

Attributes
obtained from
the same
individual

height
weight



Sequential execution, process P1

$$BMI = \frac{weight}{height^2}$$



BMI




Case 2 - computing the average of a large sample, parallel scheme

(sampled from the Brazilian male population, 1977 [3])

Attributes
obtained from
the same
individual

height
weight



$$BMI = \frac{weight}{height^2}$$



BMI




Case 2 - computing the average of a large sample, statistical scheme

(sampled from the Brazilian male population, 1977 [3])

Attributes
obtained from
the same
individual

height
weight



Sequential execution, process P1

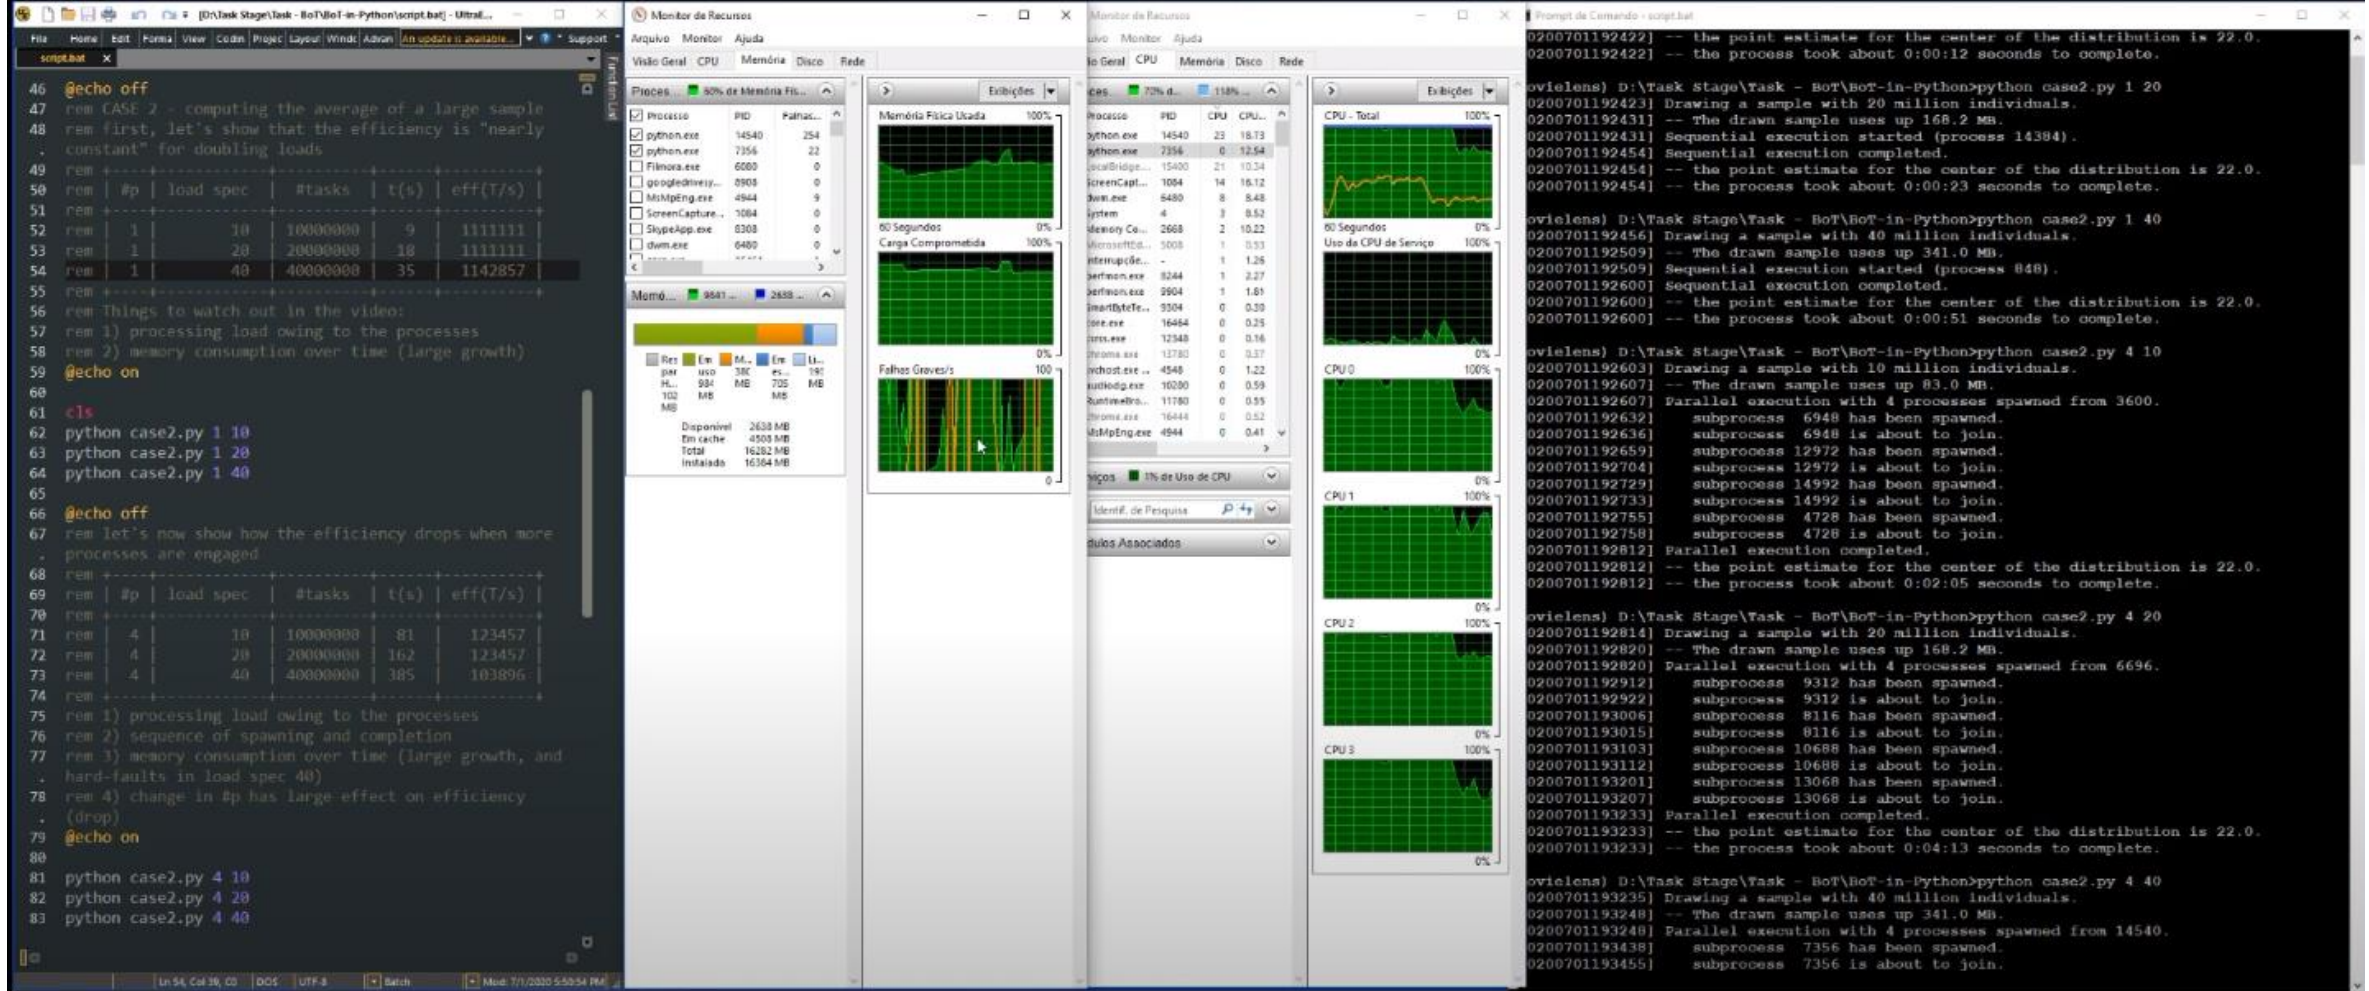
$$BMI = \frac{weight}{height^2}$$



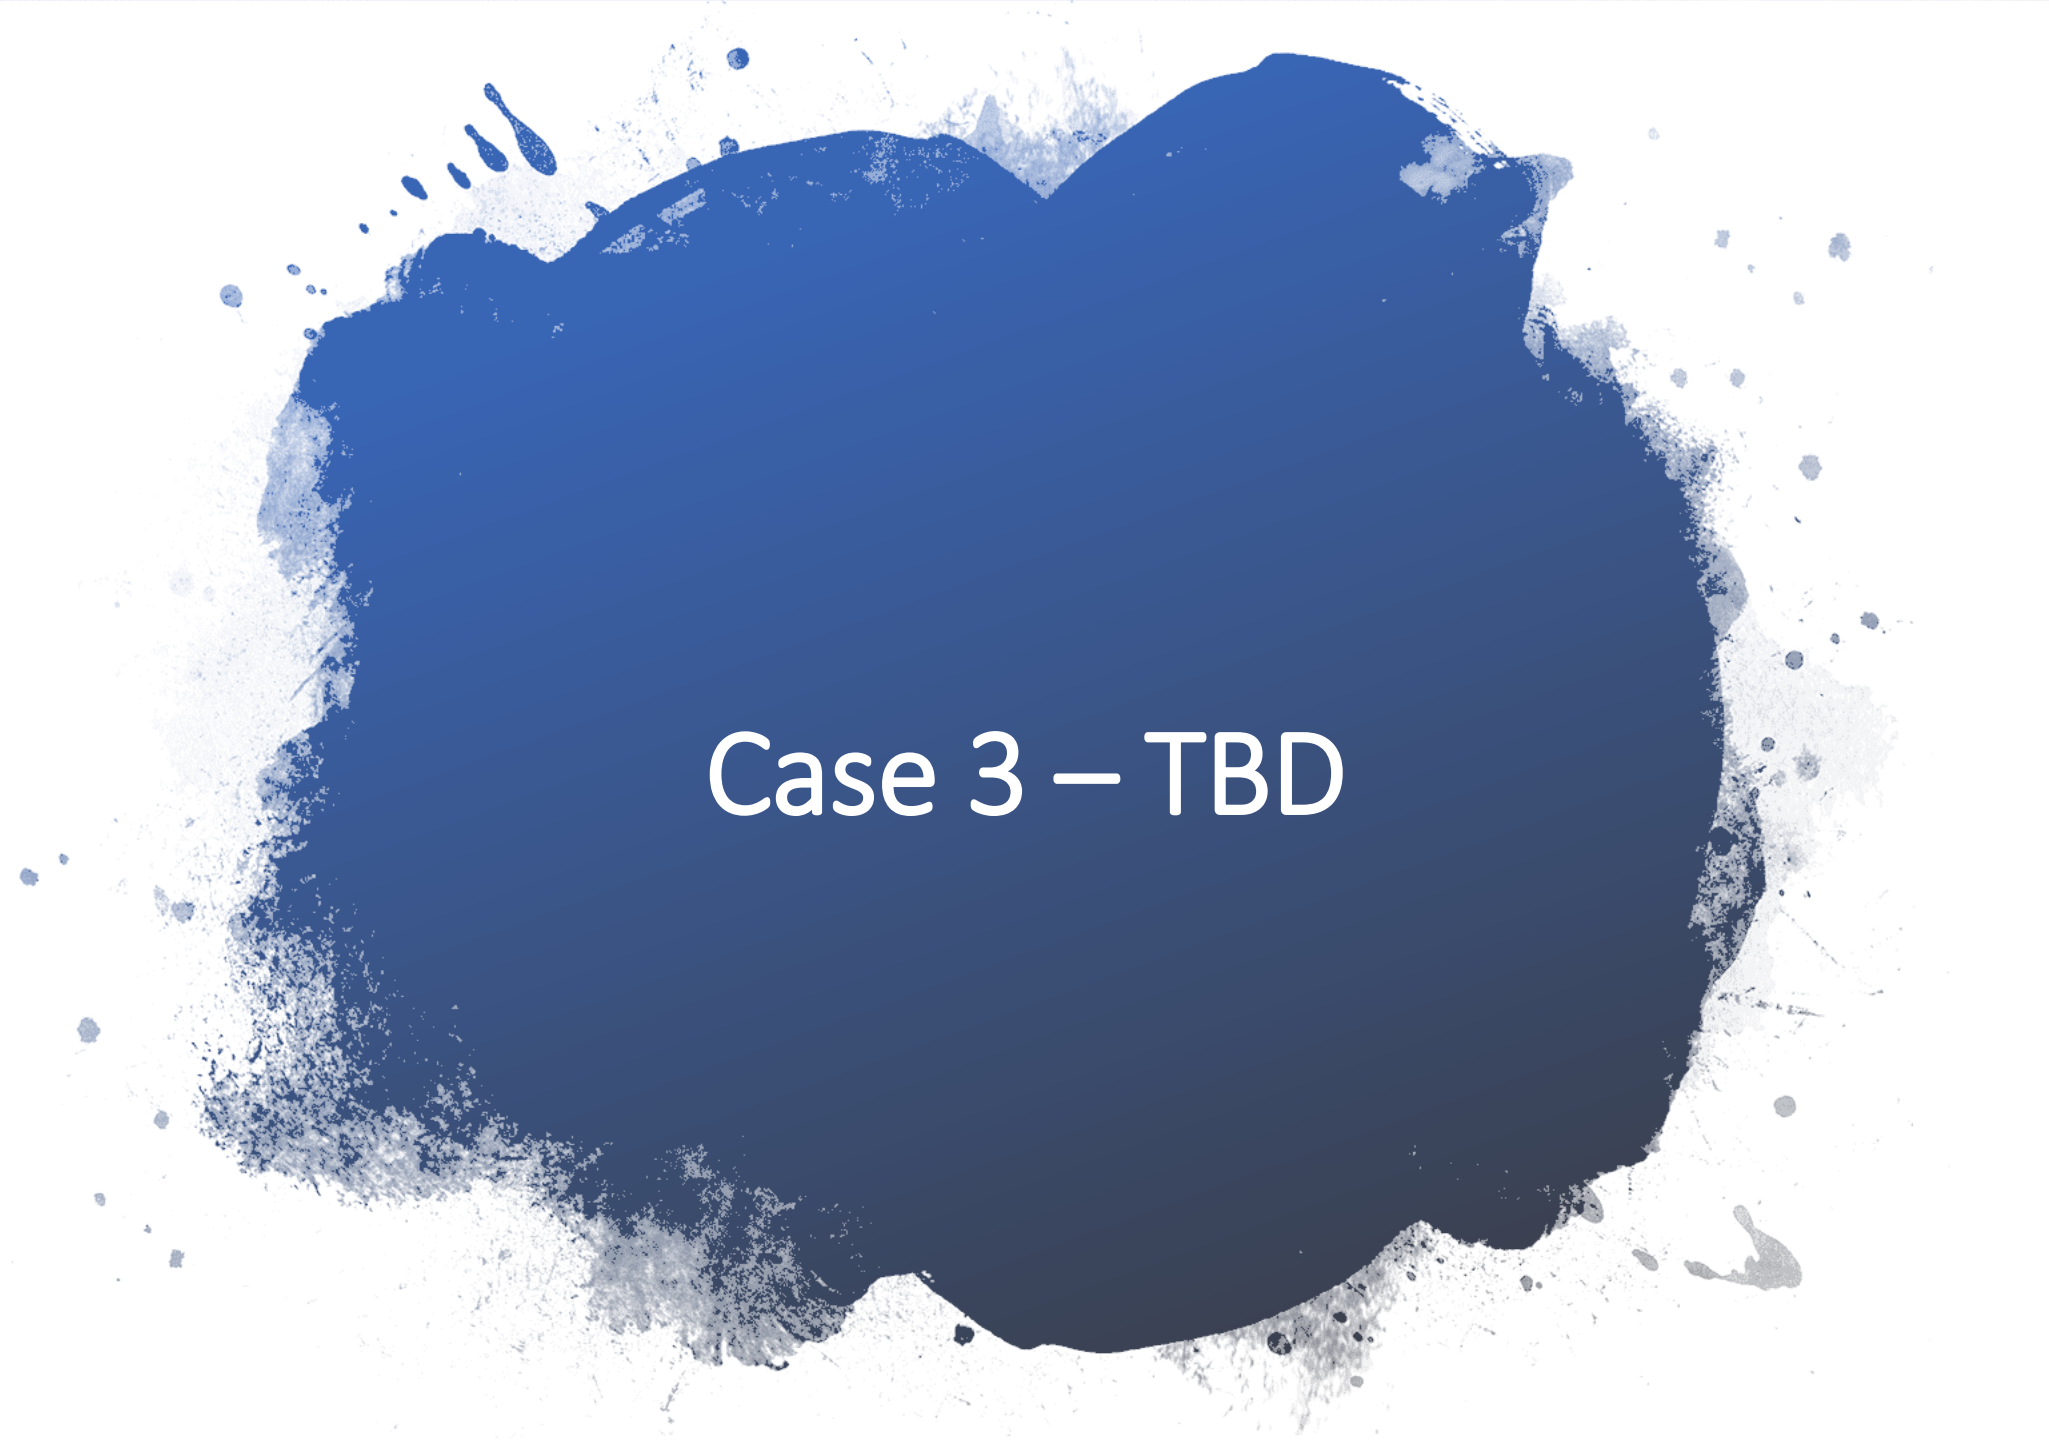
BMI



Case 2 - computing the average of a large sample



Link para o vídeo



Case 3 – TBD

To be resumed in a future meeting ;)

Parallel Programming in Python – a very, very short introduction

Thanks!

Parallel Programming in Python – a very, very short introduction

References:

[1] Quinn, Michael. "Parallel Programming in C with MPI and OpenMP", McGraw-Hill Science, ISBN 9780072822564, 2003.

[2] de Souza, Jaime Freire, Hermes Senger, and Fabricio AB Silva. "Escalabilidade de Aplicações Bag-of-Tasks em Plataformas Heterogêneas." In Anais Principais do XXXVII Simpósio Brasileiro de Redes de Computadores e Sistemas Distribuídos, pp. 664-677. SBC, 2019.

[3] Guimaraes, M. I. C. C., & Sordi, G. M. A. A. (1995). *Desenvolvimento do manequim matemático do homem brasileiro para cálculos de dosimetria interna*. Tese de Doutorado. Instituto de Pesquisas Energéticas e Nucleares (IPEN) - Universidade de São Paulo, São Paulo.
(see page 29; uses data that were collected by IBGE back in 1976/1977)